## Not Implemented Due to the Time Limit

Given the assignment time limit, I focused on a small working pipeline instead of adding more complex NLP components.

If more time were available, I would add:

- A pretrained clinical NLP model and/or LLM for concepts such as negation, uncertainty, and temporality.
- Multiple extraction methods and compare their outputs.
- A controlled majority/agreement approach between methods instead of relying on one model.
- A larger manually reviewed evaluation set.
- More extensive clinical terminology mapping.
- More organization-specific configurations and validation.

For clinically important concepts such as diagnoses, I would also consider **risk-based weighting**. For example, false negatives for a critical diagnosis may have a higher cost than false positives, so the review threshold could be stricter for those concepts.

This weighting would affect the decision and review policy, not the underlying extraction confidence score.

## Tool Disclosure

**I used AI assistance during development for code organization, debugging, and converting my Jupyter Notebook into suitable project files due to the limited time available for the assignment.**